



3000788 Intro to Comp Molec Biol

Week 5: RNA sequencing

Fall 2024



Sira Sriswasdi, PhD

- Research Affairs
- Center of Excellence in Computational Molecular Biology (CMB)
- Center for Artificial Intelligence in Medicine (CU-AIM)

Part I: RNA sequencing & differential expression




- Which class of RNA molecules are you interested in?
- Gene level or isoform level
- Do you want to discover new isoform?
- How to quantify gene expression?
- DESeq2 (read count) vs sleuth (TPM) pipeline




RNA-seq scopes


Non-coding RNAs



Total RNA sequencing





Binding of ssDNA Probes




Single-stranded DNA probes hybridize specifically to rRNA molecules.

rRNA Degradation by Ribonuclease H (RNase H) Enzyme



RNase H degrades the hybridized RNA (rRNA).

Probe Degradation by DNase I Enzyme & Clean Up



DNase I degrades the DNA probes.

rRNA-depleted RNA




Non-rRNA species (blue) are enriched.

Source: New England BioLabs



mRNA and miRNA sequencing

Isolate Total RNA

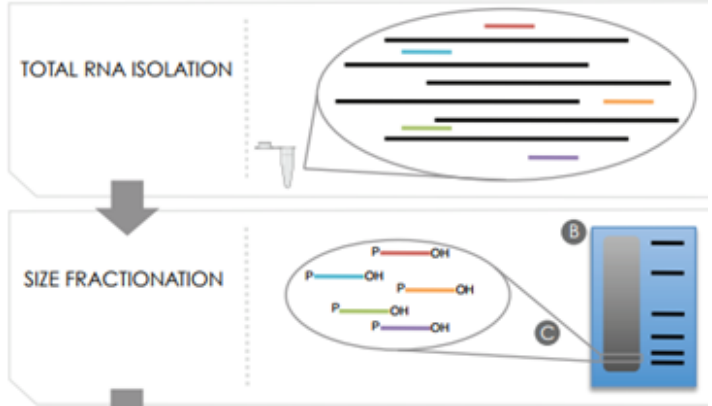



Fragmentation
and/or Isolation

In this case, isolation via Poly(T)
coated magnetic beads



Poly(A) RNA molecules
bind to the Poly(T)
magnetic beads



- Selection by polyT probe
- Size fractionation

Full-length transcript sequencing



Choosing RNA-seq technique




- miRNA or mRNA or total-RNA
- Single-end or paired-end
 - Single-end is ok for gene-level quantification
 - Paired-end is needed to distinguish isoforms
- Illumina or 3rd generation sequencer
 - Long-read data is helpful for genes with complex isoforms and for detecting novel isoforms



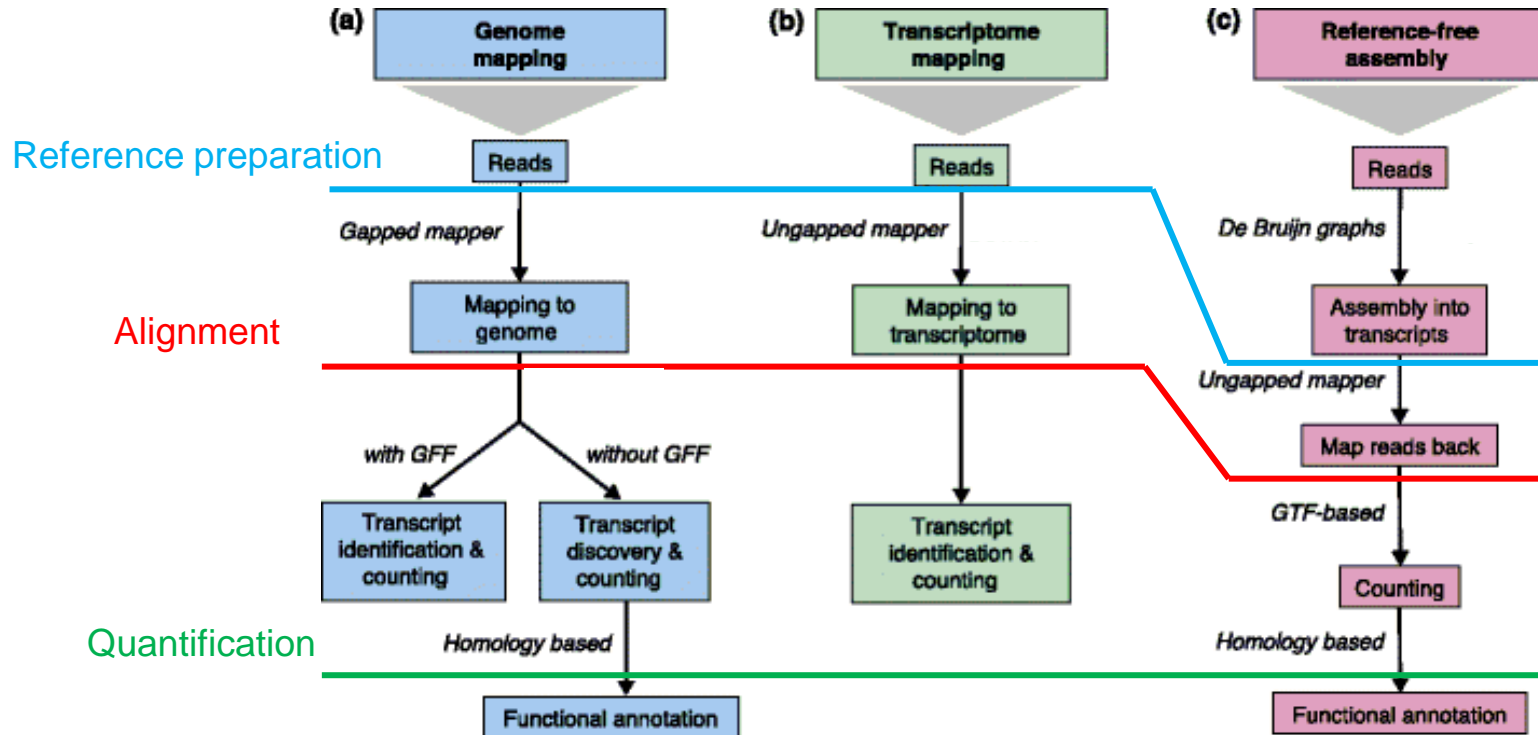
RNA-seq analysis

Three primary pipelines

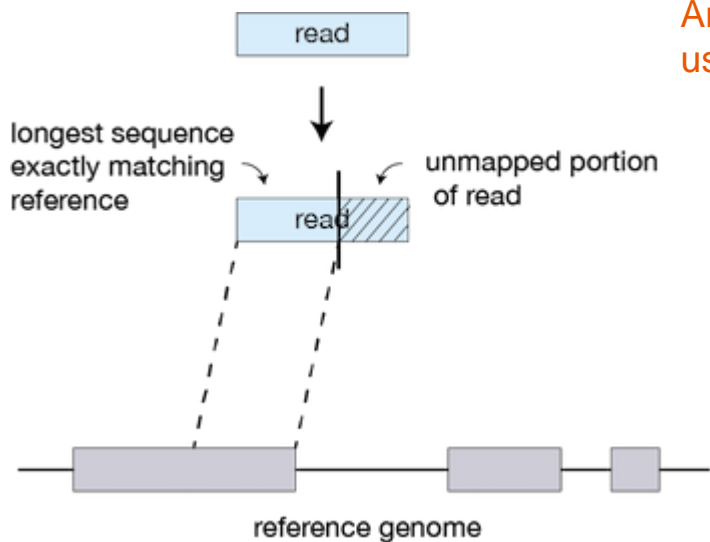
- Reference-free
 - Novel species, rely on *de novo* assembly
- Reference transcriptome
 - **Fast**, cannot discover new isoform
 - Ungapped, *k*-mer-based alignment
- Reference genome
 - **Slow**, but can detect new isoforms
 - Gapped alignment




Pipeline overview



Gapped alignment



Annotated exon structure can be used to support the alignment



GTF/GFF genome annotation format



Sample GTF output from Ensembl data dump:

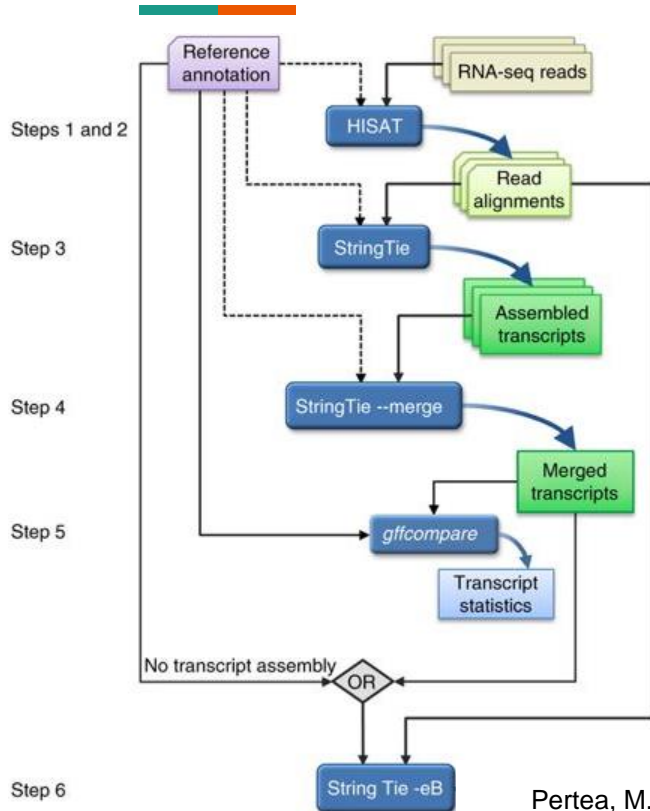
```
1 transcribed_unprocessed_pseudogene   gene      11869 14409 . + . gene_id "ENSG00000223972"; gene_name "DDX11L1"; gene_source "havana";
1 processed_transcript                  transcript 11869 14409 . + . gene_id "ENSG00000223972"; transcript_id "ENST00000456328"; gene_name
```

Sample GFF output from Ensembl export:

X	Ensembl Repeat	2419108	2419128	42	.	.	hid=trf; hstart=1; hend=21
X	Ensembl Repeat	2419108	2419410	2502	-	.	hid=AluSx; hstart=1; hend=303
X	Ensembl Repeat	2419108	2419128	0	.	.	hid=dust; hstart=2419108; hend=2419128
X	Ensembl Pred.trans.	2416676	2418760	450.19	-	2	genscan=GENSCAN00000019335
X	Ensembl Variation	2413425	2413425	.	+	.	
X	Ensembl Variation	2413805	2413805	.	+	.	


- Tab-separated text file
- Chromosome ID, object name, base pair positions, strand, and other annotation details

Multi-step post-alignment processing



- Initial alignment
- Assemble potential novel isoforms
- Merge isoforms across samples
- Re-quantify isoform abundances using the merged database of isoforms

Importance of merging isoforms



Pertea, M. et al. Nature Protocols 11:1650-1667 (2016)

- Rare isoform may be missing in some samples
- Reads can get misinterpreted if the correct isoform is not in the reference

GTF with abundance annotation



chr1	StringTie	transcript	36534	36849	1000	.	.	gene_id "STRG.1"; transcript_id "STRG.1.1"; cov "19.614035"; FPKM "6.688056"; TPM "10.944590";
chr1	StringTie	transcript	35245	36073	1000	-	.	gene_id "STRG.2"; transcript_id "STRG.2.1"; reference_id "ENST00000461467.1"; ref_gene_id "ENSG00000237613.2";
								ref_gene_name "FAM138A"; cov "0.327684"; FPKM "0.111735"; TPM "0.182847";
chr1	StringTie	transcript	52473	53312	1000	+	.	gene_id "STRG.3"; transcript_id "STRG.3.1"; reference_id "ENST00000606857.1"; ref_gene_id "ENSG00000268020.3";
								ref_gene_name "OR4G4P"; cov "0.119048"; FPKM "0.040593"; TPM "0.066429";
chr1	StringTie	transcript	137682	137965	1000	-	.	gene_id "STRG.4"; transcript_id "STRG.4.1"; reference_id "ENST00000595919.1"; ref_gene_id "ENSG00000269981.1";
								ref_gene_name "RP11-34P13.16"; cov "0.000000"; FPKM "0.000000"; TPM "0.000000";
chr1	StringTie	transcript	139283	139642	1000	.	.	gene_id "STRG.5"; transcript_id "STRG.5.1"; cov "3.111111"; FPKM "1.060837"; TPM "1.735993";

- Different tool outputs transcript abundance in different format
- GTF can accommodate abundance annotation in the last columns
 - Coverage (cov) = fraction of transcript length with mapped read
 - FPKM = **F**ragment **p**er **k**ilobase of exon per **m**illion reads mapped
 - TPM = **T**ranscript **p**er **m**illion

Units for transcript abundance

$$\text{FPKM} = \frac{\text{Read Count}}{\frac{\text{Transcript Length}}{1,000} \times \frac{\text{Total Read Count}}{1,000,000}}$$

Long transcript generates more fragments and more read counts

Experiment with higher sequencing depth generates more read counts

Similar to percentage (but per million)

$$\text{TPM} = \frac{\text{FPKM}}{\sum \text{FPKM}} \times 1,000,000$$

- Read count (number of mapped reads)
- FPKM = **F**ragment **p**er **k**ilobase of exon per **m**illion reads mapped
- TPM = **T**ranscript **p**er **m**illion

Alignment-based pipeline summary




- Initial alignment to reference genome (with annotated gene structure)
 - STAR / HISAT2
- [Optional]
 - Identify novel isoforms
 - Merge isoforms across samples
- Quantify transcript abundances
 - Read count / FPKM / TPM
 - StringTie2 / htseq-count



k-mer pseudoalignment


k -mer database for transcriptome



Bray *et al.* Nat Biotech 34:525-527 (2016)

- Create de Bruijn graph with k -mer as nodes
- Map node to transcripts with that k -mer
- Contig = a path on de Bruijn graph that mapped to the same transcript


k-mer pseudoalignment



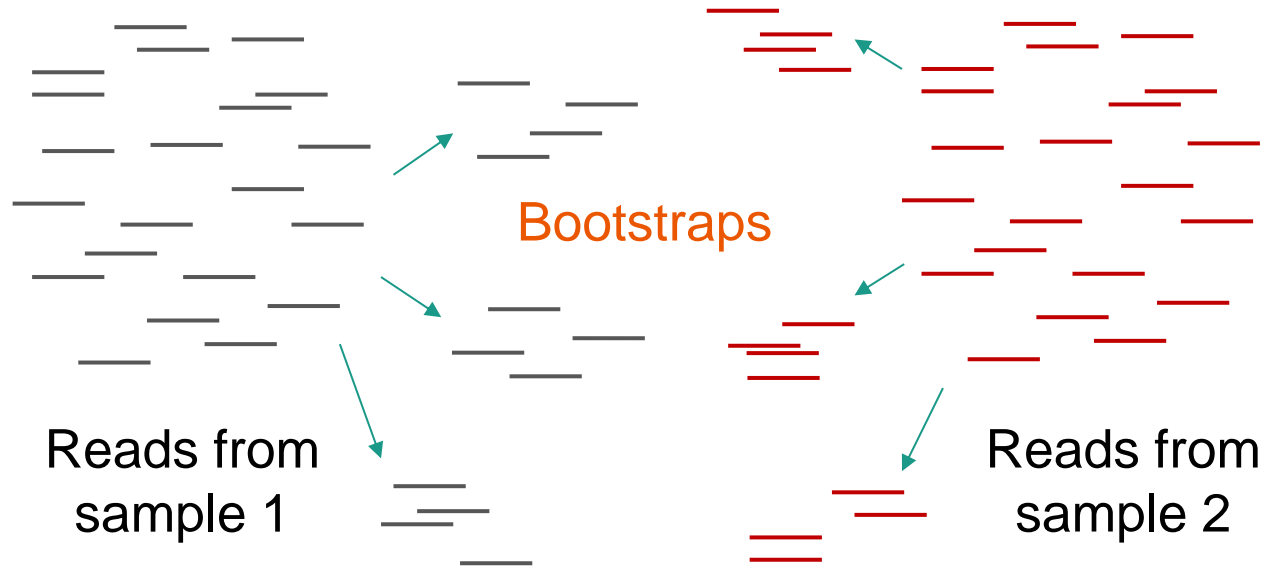
Bray *et al.* Nat Biotech 34:525-527 (2016)

- For a new read, all of its *k*-mers are mapped against contigs
 - Ignore the ordering of *k*-mers on the read
- Report only contigs that are compatible with all *k*-mers
- Speed up by skipping uninformative *k*-mer
 - V_2 and V_3 regions
 - Only 2-4 *k*-mer lookups are enough

>100 fold speed up with pseudoalignment




Bootstrapping enabled by pseudoalignment



- Bootstrapping estimates technical variances

Salmon: improved k -mer alignment

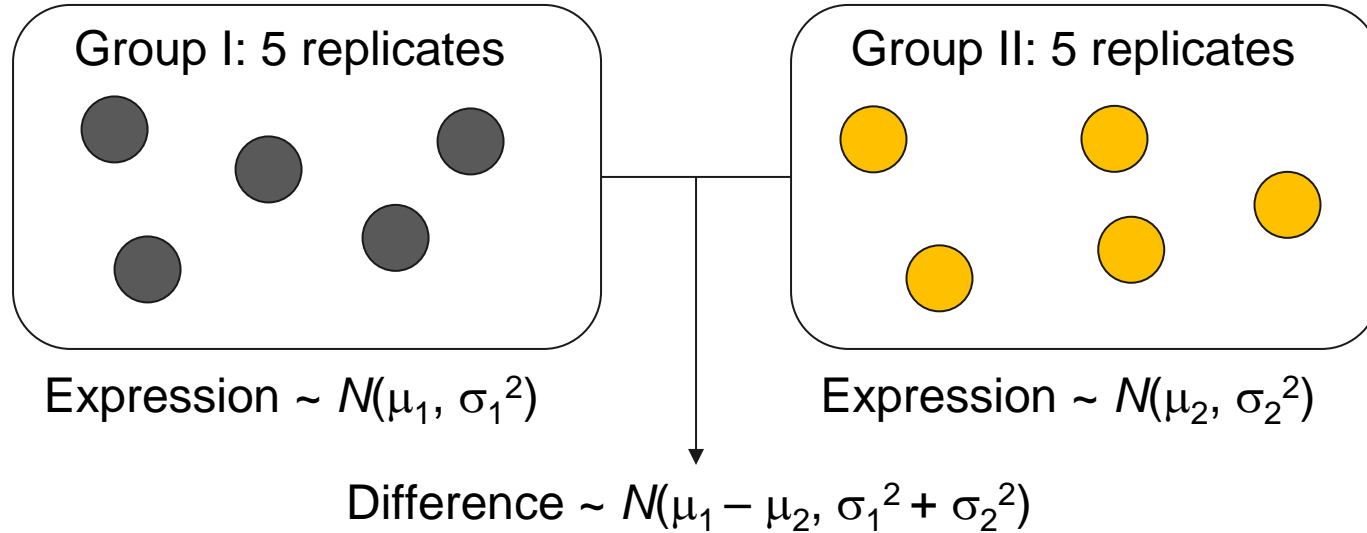


- Also track the location of k -mer on the input read → semi-alignment
- Correct quantification based on GC content and 3' / 5' amplification biases



Differential expression analysis

DE with t -test



- Simple t -test for normally distributed abundance data

DE as nested model testing / likelihood ratio test



Hypothesis 1




Hypothesis 2



Hypothesis 3



DE as permutation test




Tyanova and Cox. Cancer Systems Biology pp 133-148. (2018)

- Permuting sample labels = remove condition-specificity



DESeq2 model for read count

The distribution of RNA-seq read count



Negative binomial model

- $NB(r, p)$ = the number of failures that we will see in a series of Bernoulli trials with probability of success p until we obtain r successes
 - $X O O X X X O O X O O$ = 5 failures until 6 successes
- $P_{NB}(k; r, p) = \binom{k+r-1}{k} (1-p)^k p^r$
 - $k + r - 1$ locations to place k failures (the last location must be success)
- Mean = $\frac{pr}{(1-p)}$
- Variance = $\frac{pr}{(1-p)^2} = \frac{pr}{(1-p)} \frac{1}{(1-p)} = \frac{pr}{(1-p)} \left(1 + \frac{p}{1-p}\right)$
$$= \frac{pr}{(1-p)} + \frac{p^2 r}{(1-p)^2} = \frac{pr}{(1-p)} + \left(\frac{pr}{(1-p)}\right)^2 \frac{1}{r}$$

Another view of negative binomial model



- $P_{\text{NB}}(k; r, p) = \int_0^\infty P_{\text{Poisson}(\lambda)}(k) \cdot P_{\text{Gamma}\left(r, \frac{1-p}{p}\right)}(\lambda) d\lambda$
- Negative binomial distribution is a continuous mixture of Poisson distribution, with mixing weights Gamma-distributed
 - Same as Gamma site-specific mutation rates
- Bulk gene expression is an average over many cells
- Mixture of read counts from multiple cells, each following $\text{Poisson}(\lambda)$

DESeq2 model of gene expression




- Read count $K_{i,j} \sim \text{NB}(\mu_{i,j}, \sigma_{i,j}^2)$, for gene i from sample j
 - $\mu_{i,j}$ = sample effects x sequencing effects
 - $\sigma_{i,j}^2 = \mu_{i,j} + \text{gene-specific effects} \times \mu_{i,j}^2$
- Sample effects
 - Control / Treatment
 - Confounding factors: age, time after treatment, etc.
 - Log FC = $\sum_r x_{j,r} \beta_{i,r}$ where $x_{j,r}$ are design parameters for sample j and $\beta_{i,r}$ are the effect sizes
 - Linear effect model

Linear effect models

Covariates: quantitative measurements (e.g. age)

Regression model


$$\text{expression} = \beta_0 + \beta_1 \text{age}$$



Factors: categorical variables (e.g. genotype)


Means model

$$\text{expression} = \beta_1 \text{wildtype} + \beta_2 \text{mutant}$$



Mean-reference model

$$\text{expression} = \beta_1 + \beta_2 \text{mutant}$$



Legend

- Original data points

— Expected gene expression
(based on model)

- - - Expected gene expression
(of non-reference levels in mean-reference model)

Law, C.E. et al. F100Res 9:1444 (2020)

Linear model for multiple effects

Model

$$E(y) = 1.03 + 1.09x_1 + 1.97x_2 + 0.82x_1x_2$$

$$E(y) = 1.03 = 1.03 \quad (\text{for control})$$

$$E(y) = 1.03 + 1.09 = 2.12 \quad (\text{for treatment I})$$

$$E(y) = 1.03 + 1.97 = 3.00 \quad (\text{for treatment II})$$

$$E(y) = 1.03 + 1.09 + 1.97 + 0.82 = 4.90 \quad (\text{for treatments I \& II})$$

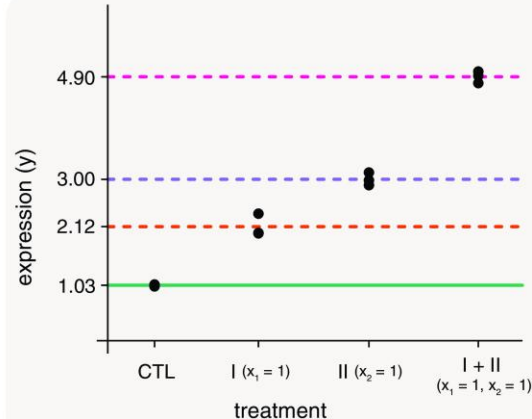
Matrix

```
> model.matrix(~treat1 * treat2)
```

$$\begin{matrix} & \text{(Intercept)} & \text{treat1YES} & \text{treat2YES} & \text{treat1YES:} \\ & & & & \text{treat2YES} \\ \begin{matrix} 1 \\ 2 \\ 3 \\ 4 \\ 5 \\ 6 \\ 7 \\ 8 \\ 9 \\ 10 \\ 11 \\ 12 \end{matrix} & \begin{pmatrix} 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 \\ 1 & 1 & 0 & 0 \\ 1 & 1 & 0 & 0 \\ 1 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 \\ 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \end{pmatrix} \end{matrix}$$

Plot

Law, C.E. et al. F100Res 9:1444 (2020)



DESeq2 model of gene expression




- Read count $K_{i,j} \sim \text{NB}(\mu_{i,j}, \sigma_{i,j}^2)$, for gene i from sample j
 - $\mu_{i,j}$ = sample effects x sequencing effects
 - $\sigma_{i,j}^2 = \mu_{i,j} + \text{gene-specific effects} \times \mu_{i,j}^2$
- Sequencing effects
 - Sequencing depth (sample-specific)
 - GC content (gene-specific)
 - Gene length (gene-specific)

DESeq2 model of gene expression



- Read count $K_{i,j} \sim \text{NB}(\mu_{i,j}, \sigma_{i,j}^2)$, for gene i from sample j
 - $\mu_{i,j}$ = sample effects x sequencing effects
 - $\sigma_{i,j}^2 = \mu_{i,j} + \text{gene-specific effects} \times \mu_{i,j}^2$
- Gene-specific effects on variance
 - **Assumption:** Genes with similar expression should have similar variances
 - Regression of **gene-specific effects** versus $\mu_{i,j}$
 - Also called **dispersion**

Two-step Bayesian approach for dispersion fitting



- **Dispersion** = $\frac{\sigma_{i,j}^2 - \mu_{i,j}}{\mu_{i,j}^2} = \left(\frac{\sigma_{i,j}}{\mu_{i,j}}\right)^2 - \frac{1}{\mu_{i,j}}$
- For genes with high expression level,
 $\text{Log}(\text{Dispersion}) \approx 2 \cdot \text{Log}\left(\frac{\sigma_{i,j}}{\mu_{i,j}}\right)$
- Fit trend using local regression
 - Similar to moving average

DE as a test of effect size



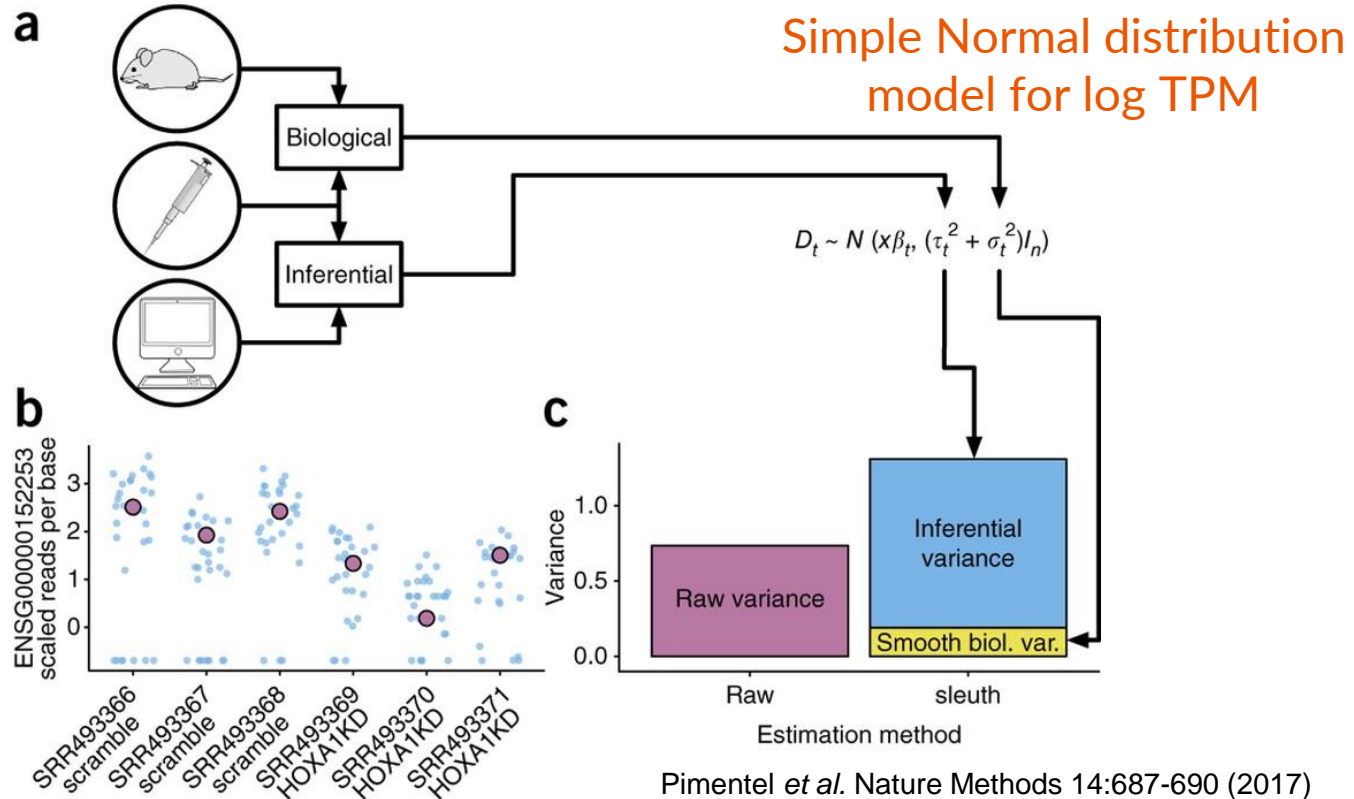
- Sample effects
 - $\text{Log FC} = \sum_r x_{j,r} \beta_{i,r}$ where $x_{j,r}$ are design parameters for sample j and $\beta_{i,r}$ are the effect sizes
- Wald test for each $\beta_{i,r}$: $\frac{\beta_{i,r}}{\text{SE}(\beta_{i,r})} \sim \text{Standard Normal}$



sleuth model for TPM

Making use of bootstrap to estimate variance

Technical
variance
estimates from
bootstrapping



Normal distribution model for log TPM




- True expression: $y_{t,i} = x_i^T \beta_t + \varepsilon_{t,i}$ for sample i and transcript t
- Observed expression: $D_{t,i} = y_{t,i} + \zeta_{t,i}$
- Noises are normally distributed: $\varepsilon_{t,i} \sim N(0, \sigma_t^2)$ and $\zeta_{t,i} \sim N(0, \tau_t^2)$
 - Transcript-specific
- Full model: $D_t \sim N(x^T \beta_t, (\sigma_t^2 + \tau_t^2) I_n)$


Variance estimates improve sensitivity of DE



Gene-level analysis



Transcript-level analysis



Differential expression summary



- DE can be formulated in multiple ways but depend heavily on the model of gene expression distribution
- **Read count** model using Negative Binomial distribution
 - Bayesian update to improve the estimate of variance
 - Tied to genome-based pipeline: **STAR**
- Log **TPM** model using Normal distribution
 - Estimate technical variance directly using bootstrapping
 - Tied to transcriptome-based pipeline with *k*-mer pseudoalignment
 - **kallisto** / **salmon**

Any question?




Part II: Functional enrichment analysis



- From individual genes to sets of genes with common characteristics
- Overrepresentation = Fisher's exact test = hypergeometric distribution
- Gene Set Enrichment Analysis (GSEA)
- Network topology-based

Differential expression result



- Statistical significance (p-values) and effect size (fold-changes)
- Do these genes correspond to specific biological characteristics?



Overrepresentation analysis

Enrichment fold



Gene group	Kinase	Not kinase	Total
Differentially expressed	50	350	400
Not differentially expressed	150	5450	5600
Total	200	5800	6000

- There are 200 kinases among 6000 genes
- Expected $400 \times 200 / 6000 = 13$ kinases to be differentially expressed
- Enrichment = $50 / 13 = 3.85$ folds


Fisher's Exact Test



Gene group	Kinase	Not kinase	Total
Differentially expressed	$k \geq 50$	$400 - k$	400
Not differentially expressed	$200 - k$	$5400 + k$	5600
Total	200	5800	6000

- P-value for this observation = $P(\text{Kinase} \ \& \ \text{DE} \geq 50)$
- $P(\text{Kinase} \ \& \ \text{DE} = k) = \text{Hypergeometric}(N = 6000, K = 200, n = 400, k)$

Hypergeometric distribution




- $\binom{K}{k}$ ways to select the intersected k genes
- $\binom{N-K}{n-k}$ ways to select the remaining $n - k$ non-kinase genes
- Total of $\binom{N}{n}$ ways
- Probability = $\frac{\binom{K}{k} \binom{N-K}{n-k}}{\binom{N}{n}}$



Gene Set Enrichment Analysis (GSEA)


GSEA algorithm sketch

A Phenotype
Classes
A B




B

Gene set




- Sort genes by the extent of up-/down-regulation across conditions
- Label genes annotated with a function
- If these genes are clustered together at the **top**, then this function is **up-regulated**
- If these genes are clustered together at the **bottom**, then this function is **down-regulated**

GSEA scoring




- Starting at **score** = 0 from the top of the sorted gene list
- If encounter gene from S , +**score**
- Otherwise, -**score**
- **Score** indicates the extent of up-/down-regulation
 - Correlation with conditions
 - Log fold-change

Up-regulated versus unchanged pathways




Null hypothesis for GSEA



- **Null hypothesis:** Genes from S are uniformly distributed in the list
- +score and -score are uniformly distributed in the list
- This is a **Random Walk**


Statistical behaviors of random walks



<https://demonstrations.wolfram.com/SimulatingTheSimpleRandomWalk/>

- $P(\text{maximal deviation} > d) \approx 2 \sum_{k=1}^{\infty} (-1)^{k-1} e^{-2(kd)^2}$

Kolmogorov-Smirnov test



- Test whether two probability distribution are equal
- Compare cumulative density (red and blue trends)
- If they are equal, the two curves should stay close to each other
- **Null hypothesis:** random walk

Setting the score for GSEA

Enrichment statistic. The exponential scaling factor of the phenotype score in enrichment score formula.

p ?

2 ▼

1


0

1.5

2

- Originally developed for microarray data
- Adapted to RNA-seq
 - Log fold-change
 - No score (simply rank genes)
- Weighted score = (score) ^{p}
 - Default: $p = 1$
 - No score: $p = 0$
 - More weights for top genes: $p > 1$

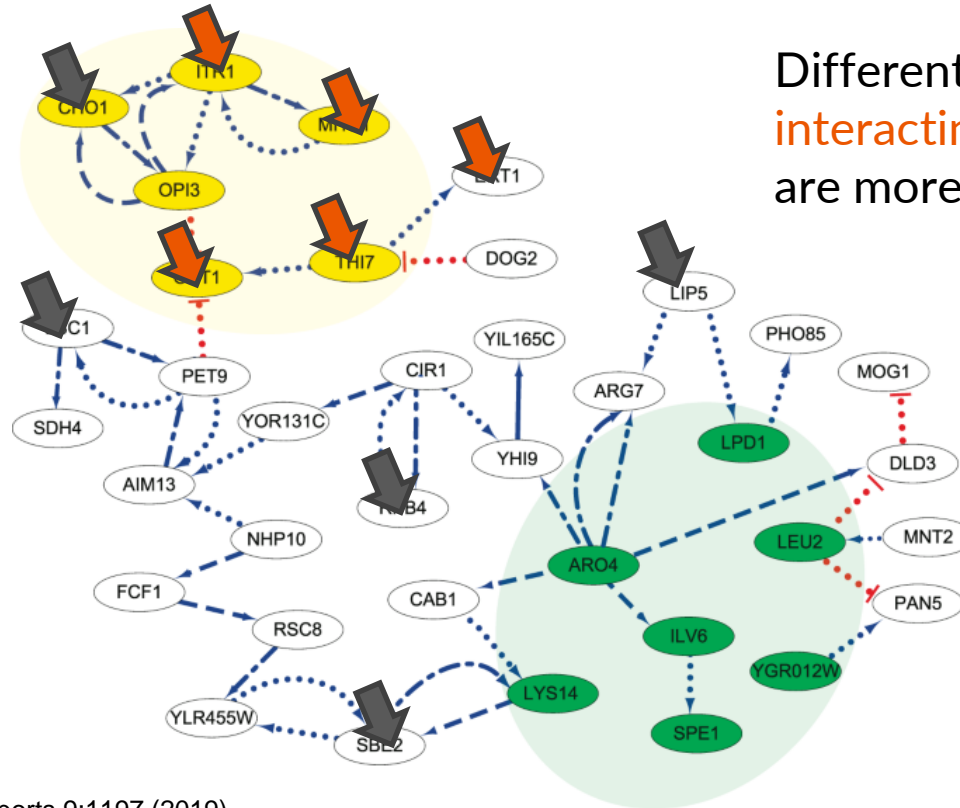
Comparing the impact of $p = 0$ and 1






Network topology-based analysis

Gene and protein interaction networks

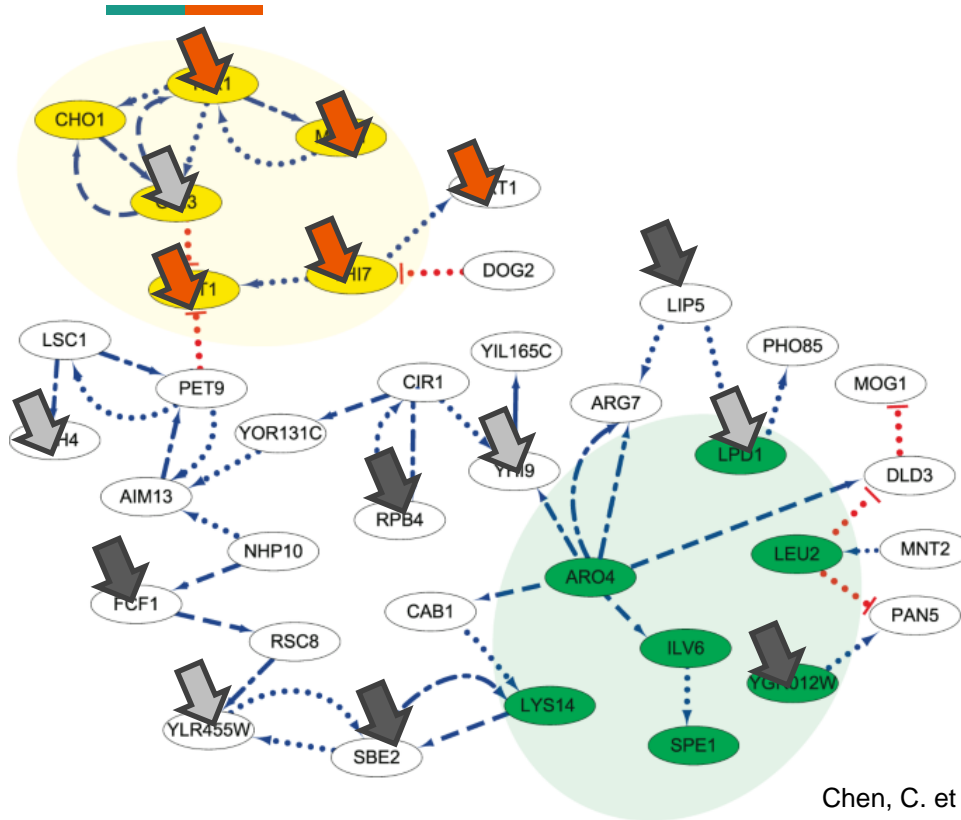


Network coherence scores



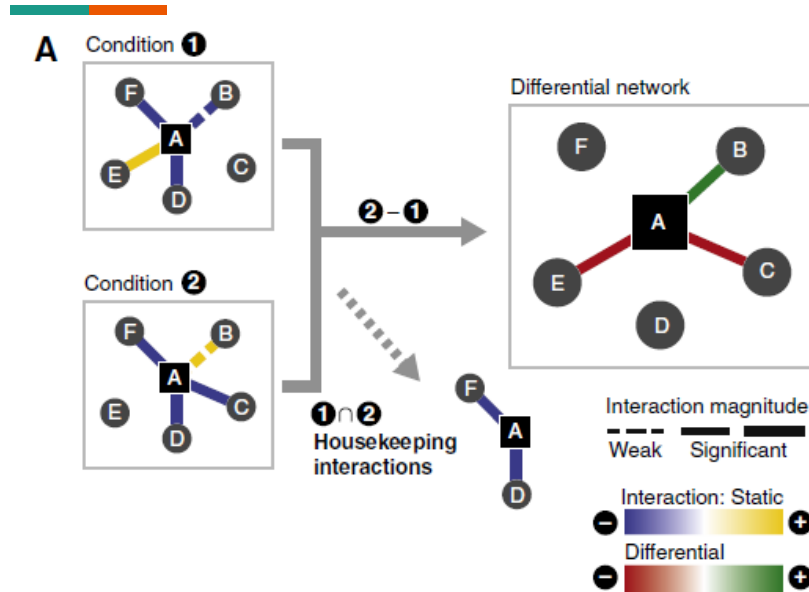
- Connectedness
 - Number of components
 - Number of edges
- Path length between genes
 - Unweighted
 - Weighted by fold changes

Permutation test: Gene set



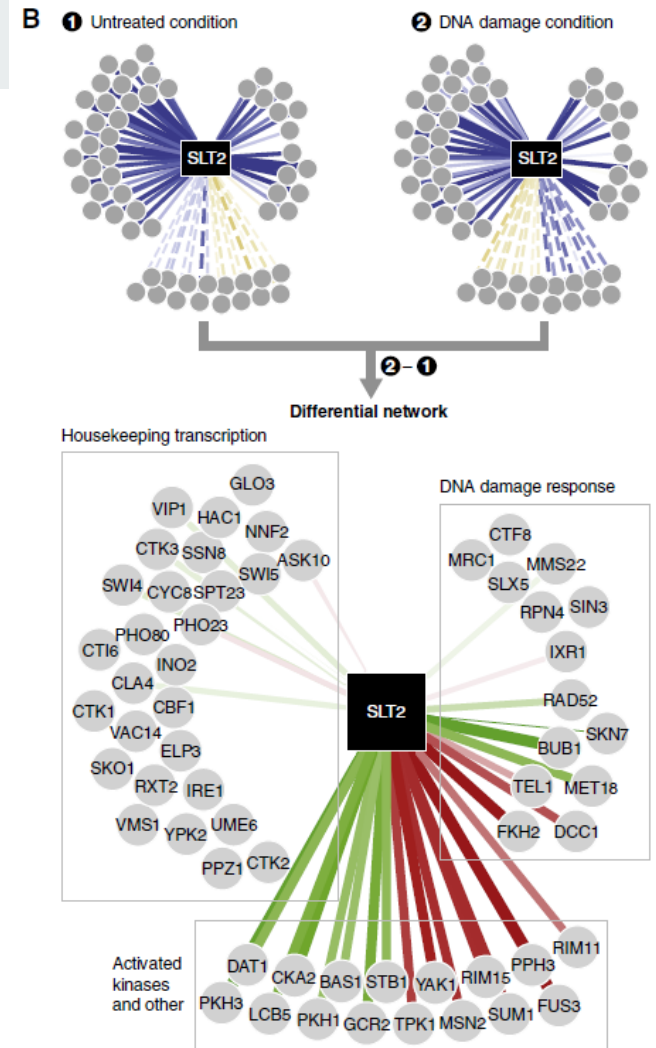
- Randomly select the same number of genes
- Recalculate network coherence scores
- P-value = fraction of samplings that the score is \geq the original

Differential network



Ideker *et al.* Mol Syst Biol, 8:565 (2012)

- Detect gain/loss gene co-expression
- Unaffected interactions remain the same



Pros and cons



- Overrepresentation
 - Easy and fast to calculate
 - Depend on p-value cutoff
- GSEA
 - No p-value cutoff
 - Distinguish up- and down-regulated functions
- Network-based
 - Most biologically meaningful
 - Network data is incomplete

Any question?

